


System Calls

1. System calls are the **user API** to the OS
2. System calls are **not** function calls!
3. **How** are system calls **invoked**?
4. **Why** are system calls implemented this way?

Operating System Interfaces: System Calls



Types of System Calls

Process Control

- load
- execute
- end, abort
- create process
- terminate process
- get/set process attributes
- wait for time, wait event, signal event
- allocate, free memory

File Management

- create file, delete file
- open, close
- read, write, reposition
- get/set file attributes

Device Management

- request device, release device
- read, write, reposition
- get/set device attributes
- logically attach or detach devices


Information Maintenance

- get/set time or date
- get/set system data
- get/set process, file, or device attributes

Communication

- create, delete communication connection
- send, receive messages
- transfer status information
- attach or detach remote devices

System Calls – They look like Functions ...



```
#include<unistd.h>
```


```
int main(void) {
```

```
    write(1, "Hello, world\n", 13);
```

```
    return 0;
```

```
}
```

... but they are not!



64-bit Linux?

- same but different
- different registers
- different numbers
- **syscall** instruction

Making System Calls in 32-bit Linux

1. load **system call number** in register **eax**.
2. load **arguments** to system call in registers **ebx, ecx, edx, esi, edi, ebp**
3. invoke **software interrupt**: **int 0x80**

Returned values are stored in **eax**.

System Call Numbers

Linux Syscall Reference

Show 10 entries Search:

#	Name	eax	ebx	ecx	edx	esi	edi	Definition
0	sys_restart_syscall	0x00	-	-	-	-	-	kernel/signal.c:2058
1	sys_exit	0x01	int error_code	-	-	-	-	kernel/exit.c:1046
2	sys_fork	0x02	struct pt_regs *	-	-	-	-	arch/alpha/kernel/entry.S:716
3	sys_read	0x03	unsigned int fd	char __user *buf	size_t count	-	-	fs/read_write.c:391
4	sys_write	0x04	unsigned int fd	const char __user *buf	size_t count	-	-	fs/read_write.c:408
5	sys_open	0x05	const char __user *filename	int flags	int mode	-	-	fs/open.c:900
6	sys_close	0x06	unsigned int fd	-	-	-	-	fs/open.c:969
7	sys_waitpid	0x07	pid_t pid	int __user *stat_addr	int options	-	-	kernel/exit.c:1771
8	sys_creat	0x08	const char __user *pathname	int mode	-	-	-	fs/open.c:933
9	sys_link	0x09	const char __user *oldname	const char __user *newname	-	-	-	fs/namel.c:2520

Showing 1 to 10 of 338 entries

First Previous 1 2 3 4 5 Next Last

Generated from Linux kernel 2.6.35.4 using **Exuberant Ctags**, **Python**, and **DataTables**.
Project on [GitHub](#). Hosted on [GitHub Pages](#).

syscalls.kernelgrok.com

System Calls are Expensive!

Software interrupts are expensive!

- Cost of context switch (saving/restoring registers)
- Caches are stale
- TLBs
- CPU pipelines

Compiler optimization not possible.

Why Interrupts or `syscall`?

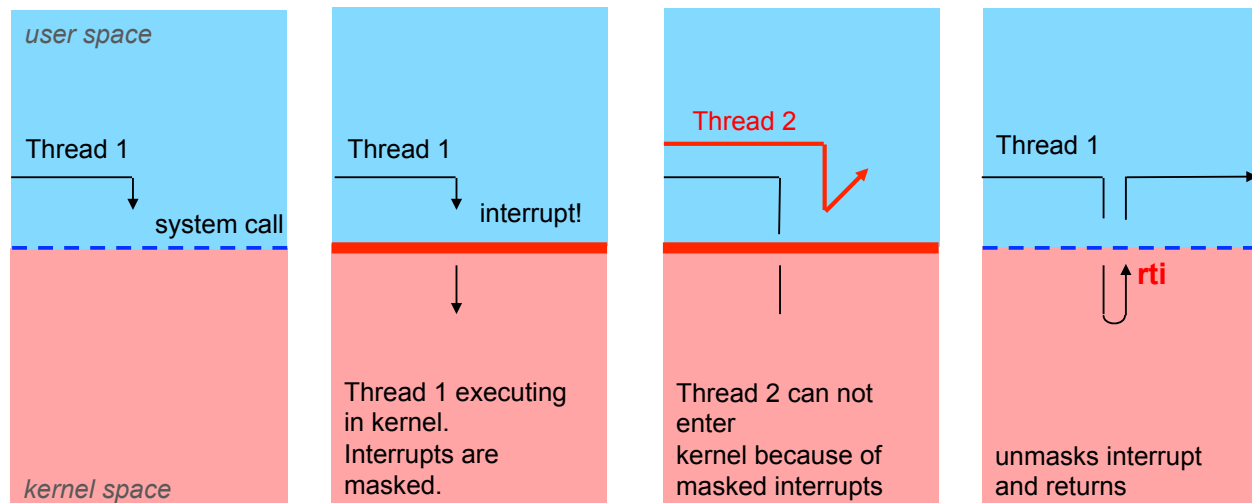
Reason 1: Can load user program into memory without knowing exact address of system functions.

Reason 2: Separation of address space, including stacks: *user stack* and *kernel stack*.

Reason 3: Automatic change to *supervisor mode*.

Reason 4: Can control *access* to kernel by masking interrupts.

Reason 4: Mutual Exclusion in Kernel



Summary: System Calls

1. System calls are the **OS API**
2. They **look like** function calls...
... but they **are not**!
3. They are implemented using **software interrupts** (or variations thereof)
4. **Why** are they implemented this way?