

Low-Level Thread Dispatching on the x86

- Dispatching
- Dispatching and Exception Handling
- RECAP: Exception Handling in the x86
- How to use Exception Handling mechanisms for Dispatching
- RECAP: The Thread Control Block (TCB) in class **Thread**
- Dispatching in 4 Steps: Start, Store, Locate, Load
- Creating a Thread
- The Execution Lifecycle of a Thread

Low-Level Thread Dispatching

Low-level Dispatch

```
Thread::dispatch_to(Thread * _thread);
```

Store state of current thread execution and load state of target thread onto CPU.

Dispatching in 3 steps:

1. Save state of current thread
2. Load state of new thread
3. Continue execution of new thread

Low-Level Thread Dispatching

Dispatching in 3 steps:

1. **Save** state of **current** thread
2. **Load** state of **new** thread
3. **Continue** execution of new thread

Problems:

- How do we **save** the state of the current thread?
- How to we **load** the state of the new thread?
- How do we **continue** execution of new thread?!

Low-Level Thread Dispatching

Problems:

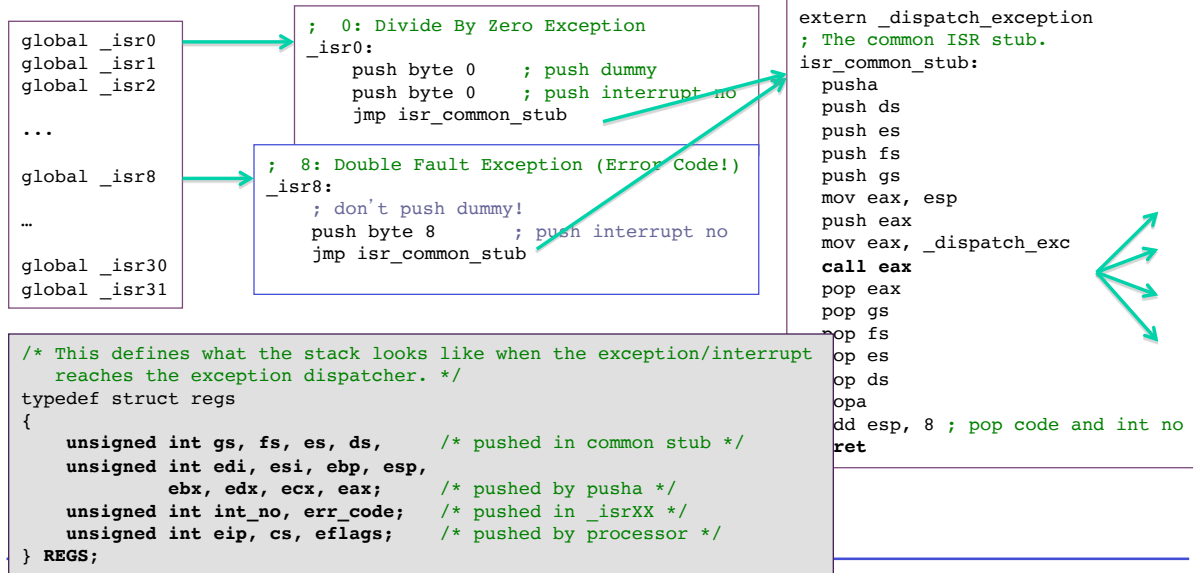
- How do we **save** the state of the current thread?
- How to we **load** the state of the new state?
- How do we **continue** execution of new thread?!

Q: Where else do we have the same problem?

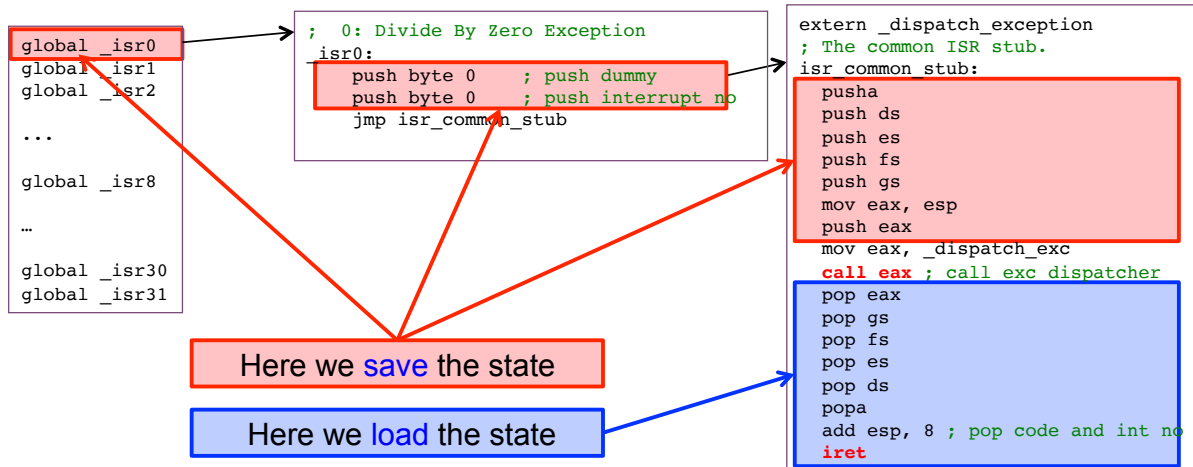
A: When we **handle exceptions**.

Solution: Handle Thread Dispatching (somewhat) in the same manner!

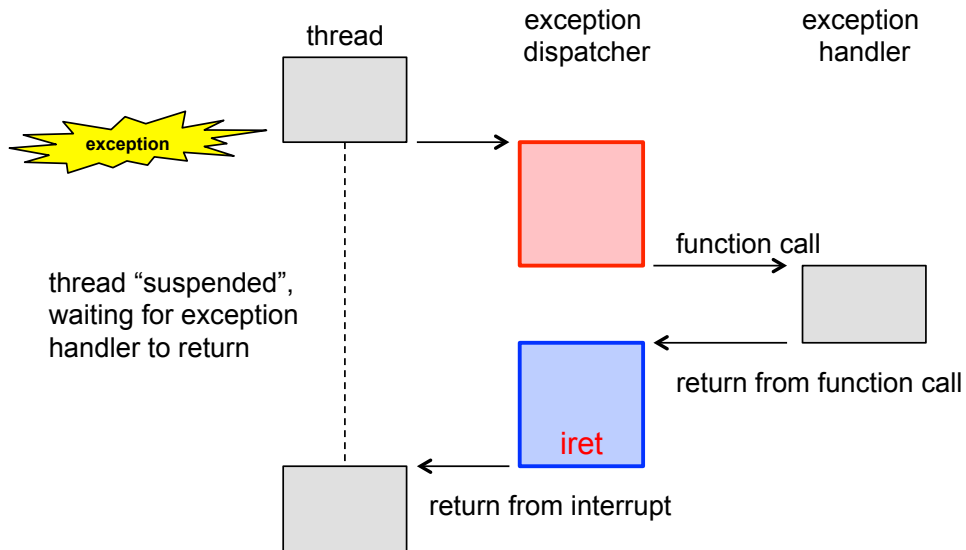
RECAP: Exception Handling




RECAP: Saving and Loading State




Saving and Restoring the State



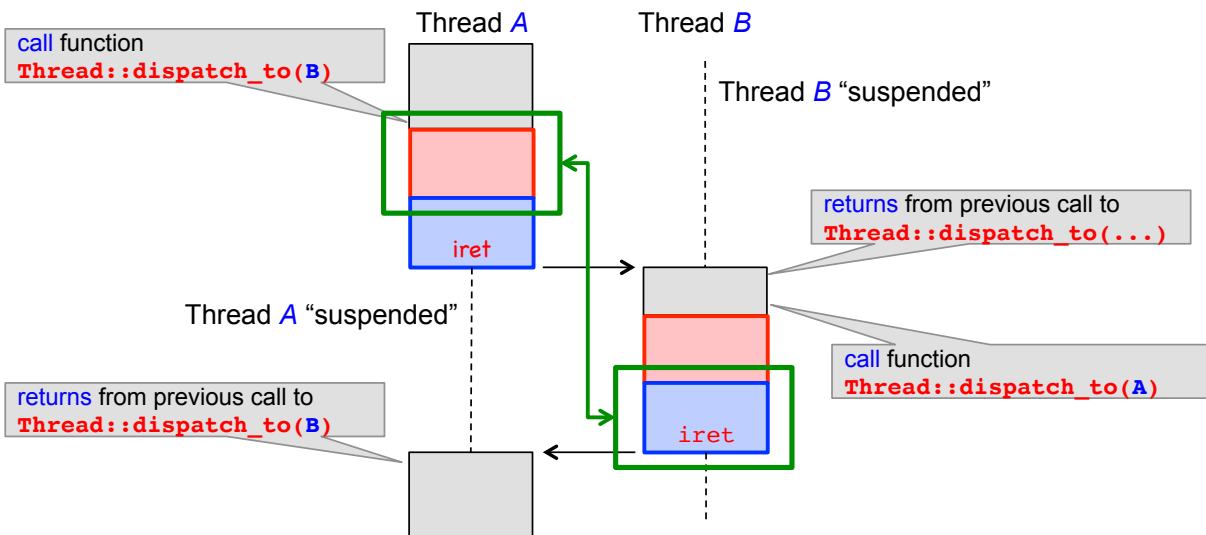
Saving and Restoring the State



Saving and Restoring the State



Dispatching, Example




RECAP: Structure of a Thread

Thread Control Block

TCB:

- thread id
- stack
- stackptr
- priority
- bookkeeping
- etc...



```
class Thread {  
private:  
    char * stack_ptr;  
    /* The current stack pointer for the thread.  
       Keep it at offset 0, since the thread  
       dispatcher relies on this location! */  
    int thread_id;  
    /* thread identifier.  
       Assigned upon creation. */  
    char * stack;  
    /* pointer to the stack of the thread.*/  
    unsigned int stack_size;  
    /* size of the stack (in byte) */  
    int priority;  
    /* Maybe the scheduler wants to use priorities. */  
    etc...  
};
```


Dispatching Step-by-Step: 0. Start

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

return_addr
_thread

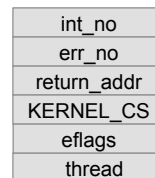
Dispatching Step-by-Step: 1. Fix Stack

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}  
_threads_low_switch_to:  
    ; STEP 1: fix stack  
    push    eax            ; save eax  
    mov     eax, [esp+4]    ; get return address  
    mov     [esp-4], eax    ; move return addr  
                           ; down 8 bytes from orig loc  
    add     esp, 8         ; move stack ptr up  
    pushfd                     ; put eflags where return address was  
    mov     eax, [esp-4]    ; restore saved value of eax  
    push    dword KERNEL_CS ; push cs selector  
    sub     esp, 4         ; point stack ptr  
                           ; at return address  
  
    ; STEP 2: ...  
    ...
```



Dispatching Step-by-Step: 2. Save State

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}  
_threads_low_switch_to:  
    ; STEP 1: fix stack  
    ...  
    ; STEP 2: save state  
    ; Push fake error code and interrupt number  
    push    dword 0  
    push    dword 0
```



Dispatching Step-by-Step: 2. Save State

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

```
    _threads_low_switch_to:  
        ; STEP 1: fix stack  
        ...  
        ; STEP 2: save state  
        ; Push fake error code and interrupt number  
        push    dword 0  
        push    dword 0  
        ; Save general purpose registers.  
        pushad
```

edi
esi
ebp
esp
ebx
edx
ecx
eax
int_no
err_no
return_addr
KERNEL_CS
eflags
_thread

Dispatching Step-by-Step: 2. Save State

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

```
    _threads_low_switch_to:  
        ; STEP 1: fix stack  
        ...  
  
        ; STEP 2: save state  
        ; Push fake error code and interrupt number  
        push    dword 0  
        push    dword 0  
        ; Save general purpose registers.  
        pushad  
        push    ds  
        push    es  
        push    fs  
        push    gs
```

gs
fs
es
ds
edi
esi
ebp
esp
ebx
edx
ecx
eax
int_no
err_no
return_addr
KERNEL_CS
eflags
_thread

Dispatching Step-by-Step: 2. Save State

```
void Thread::dispatch_to(Thread * _thread) {
    threads_low_switch_to(_thread);
}

_threads_low_switch_to:
; STEP 1: fix stack
...
; STEP 2: save state
; Push fake error code and interrupt
push    dword 0
push    dword 0
; Save general purpose registers.
pushad
push    ds
push    es
push    fs
push    gs

; Save stack pointer in the thread control block
; (at offset 0).
mov     eax, [_current_thread]
mov     [eax+0], esp
```

```
class Thread {
private:
    char * stack_ptr;
    /* The current stack pointer for the
       thread.
       Keep it at offset 0, since the thread
       dispatcher relies on this location! */
    int thread_id;
    /* thread identifier.
       Assigned upon creation. */
    char * stack;
    /* pointer to the stack of the thread.*/
    unsigned int stack_size;
    /* size of the stack (in byte) */
    int priority;
    /* Maybe the scheduler wants to use
       priorities. */
    etc...
```

gs
fs
es
ds
edi
esi
ebp
esp
ebx
edx
ecx
eax
int_no
err_no
return_addr
KERNEL_CS
eflags
_thread

Dispatching Step-by-Step: 3. Locate State

```
void Thread::dispatch_to(Thread * _thread) {
    threads_low_switch_to(_thread);
}

_threads_low_switch_to:
; STEP 1: fix stack
...
; STEP 2: save state
...

; STEP 3: locate state of new thread
mov     eax, dword [esp+INTERRUPT_STATE_SIZE]
```

INTERRUPT_STATE_SIZE (68 Bytes)

gs
fs
es
ds
edi
esi
ebp
esp
ebx
edx
ecx
eax
int_no
err_no
return_addr
KERNEL_CS
eflags
<u>_thread</u>

Dispatching Step-by-Step: 3. Locate State

```
void Thread::dispatch_to(Thread * _thread) {
    threads_low_switch_to(_thread);
}

_threads_low_switch_to:
; STEP 1: fix stack
...
; STEP 2: save state
...

; STEP 3: locate state of new thread
mov     eax, dword [esp+INTERRUPT_STATE_SIZE]
; Make the new thread current, and switch to its stack.
mov     [_current_thread], eax
mov     esp, [eax+0]

class Thread {
private:
    char * stack_ptr;
    /* The current stack pointer for the thread.
       Keep it at offset 0, since the thread
       dispatcher relies on this location! */
    etc...
} new thread
```

gs	gs
fs	fs
es	es
ds	ds
edi	edi
esi	esi
ebp	ebp
esp	esp
ebx	ebx
edx	edx
ecx	ecx
eax	eax
int_no	int_no
err_no	err_no
return_addr	return_addr
KERNEL_CS	KERNEL_CS
eflags	eflags
_somethread	<u>_thread</u>

Dispatching Step-by-Step: 4. Load State

```
void Thread::dispatch_to(Thread * _thread) {
    threads_low_switch_to(_thread);
}

_threads_low_switch_to:
; STEP 1: fix stack
...
; STEP 2: save state
...
; STEP 3: locate state of new thread
...
; STEP 4: load state of new thread
pop     gs
pop     fs
pop     es
pop     ds
```

edi
esi
ebp
esp
ebx
edx
ecx
eax
int_no
err_no
return_addr
KERNEL_CS
eflags
_somethread

Dispatching Step-by-Step: 4. Load State

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

```
    _threads_low_switch_to:  
        ; STEP 1: fix stack  
        ...  
        ; STEP 2: save state  
        ...  
        ; STEP 3: locate state of new thread  
        ...  
        ; STEP 4: load state of new thread  
        pop     gs  
        pop     fs  
        pop     es  
        pop     ds  
        popad
```

int_no
err_no
return_addr
KERNEL_CS
eflags
_somethread

Dispatching Step-by-Step: 4. Load State

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

```
    _threads_low_switch_to:  
        ; STEP 1: fix stack  
        ...  
        ; STEP 2: save state  
        ...  
        ; STEP 3: locate state of new thread  
        ...  
        ; STEP 4: load state of new thread  
        pop     gs  
        pop     fs  
        pop     es  
        pop     ds  
        popad  
        add     esp, 8        ; skip int num and error code
```

return_addr
KERNEL_CS
eflags
_somethread

Dispatching Step-by-Step: 4. Load State

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}  
  
_threads_low_switch_to:  
; STEP 1: fix stack  
...  
; STEP 2: save state  
...  
; STEP 3: locate state of new thread  
...  
; STEP 4: load state of new thread  
pop     gs  
pop     fs  
pop     es  
pop     ds  
popad  
add     esp, 8      ; skip int num and error code  
; We'll return to the place where the thread was  
; executing last.  
iret
```

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

return_addr
KERNEL_CS
eflags
_somethread

Dispatching Step-by-Step: 4. Load State

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

return_addr
KERNEL_CS
eflags
_somethread

Dispatching Step-by-Step: 4. Load State

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

_somethread

Dispatching Step-by-Step: 4. Load State

_somethread

Dispatching Step-by-Step: 4. Load State

Creating a Thread

```
void Thread::setup_context(Thread_Function _tfunction){  
    /* Address of thread shutdown function */  
    push((unsigned long) &tshutdown);  
}
```

```
inline void Thread::push(unsigned long _val) {  
    stack_ptr -= 4;  
    *((unsigned long *) stack_ptr) = _val;  
}
```

push **tshutdown**

ret_addr

Creating a Thread

```
void Thread::setup_context(Thread_Function _tfunction){
    /* Address of thread shutdown function */
    push((unsigned long) &tshutdown);
    /* Push the address of the thread function. */
    push((unsigned long) _tfunction);
}
```

push _tfunction
push tshutdown

ret_addr
ret_addr

Creating a Thread

```
void Thread::setup_context(Thread_Function _tfunction){
    /* Address of thread shutdown function */
    push((unsigned long) &tshutdown);
    /* Push the address of the thread function. */
    push((unsigned long) _tfunction);
    /* EFLAGS and instruction pointer */
    push(0);
    push(Machine::KERNEL_CS);
    push((unsigned long) &thread_start);
}
```

push thread_start
push KERNEL_CS
push 0
push _tfunction
push tshutdown

eip
es
eflags
ret_addr
ret_addr

Creating a Thread

```
void Thread::setup_context(Thread_Function _tfunction){
    /* Address of thread shutdown function */
    push((unsigned long) &tshutdown);
    /* Push the address of the thread function. */
    push((unsigned long) _tfunction);
    /* EFLAGS and instruction pointer */
    push(0);
    push(Machine::KERNEL_CS);
    push((unsigned long) &thread_start);
    /* Fake error code and interrupt number. */
    push(0); push(0);
}
```

push 0	int_no
push 0	err_code
push thread_start	eip
push KERNEL_CS	es
push 0	eflags
push _tfunction	ret_addr
push tshutdown	ret_addr

Creating a Thread

```
void Thread::setup_context(Thread_Function _tfunction){
    /* Address of thread shutdown function */
    push((unsigned long) &tshutdown);
    /* Push the address of the thread function. */
    push((unsigned long) _tfunction);
    /* EFLAGS and instruction pointer */
    push(0);
    push(Machine::KERNEL_CS);
    push((unsigned long) &thread_start);
    /* Fake error code and interrupt number. */
    push(0); push(0);
    /* Initial values for general-purpose registers. */
    push(0); ...; push(0); /* eax ... edi */
}
```

push 0	edi
push 0	esi
push 0	ebp
push 0	esp
push 0	ebx
push 0	edx
push 0	ecx
push 0	eax
push 0	int_no
push 0	err_code
push thread_start	eip
push KERNEL_CS	es
push 0	eflags
push _tfunction	ret_addr
push tshutdown	ret_addr

Creating a Thread

```
void Thread::setup_context(Thread_Function _tfunction){
    /* Address of thread shutdown function */
    push((unsigned long) &tshutdown);
    /* Push the address of the thread function. */
    push((unsigned long) _tfunction);
    /* EFLAGS and instruction pointer */
    push(0);
    push(Machine::KERNEL_CS);
    push((unsigned long) &thread_start);
    /* Fake error code and interrupt number. */
    push(0); push(0);
    /* Initial values for general-purpose registers. */
    push(0); ...; push(0); /* eax ... edi */
    /* Segment registers */
    push(Machine::KERNEL_DS); /* ds */
    push(Machine::KERNEL_DS); /* es */
    push(0); push(0); /* fs; gs */
}
```

push 0	gs
push 0	fs
push KERNEL_CS	es
push KERNEL_CS	ds
push 0	edi
push 0	esi
push 0	ebp
push 0	esp
push 0	ebx
push 0	edx
push 0	ecx
push 0	eax
push 0	int_no
push 0	err_code
push thread_start	eip
push KERNEL_CS	es
push 0	eflags
push _tfunction	ret_addr
push tshutdown	ret_addr

Execution Lifecycle of a Thread

```
void Thread::dispatch_to(Thread * _thread) {
    threads_low_switch_to(_thread);
}
```

push 0	gs
push 0	fs
push KERNEL_CS	es
push KERNEL_CS	ds
push 0	edi
push 0	esi
push 0	ebp
push 0	esp
push 0	ebx
push 0	edx
push 0	ecx
push 0	eax
push 0	int_no
push 0	err_code
push thread_start	eip
push KERNEL_CS	es
push 0	eflags
push t_fun_1	ret_addr
push tshutdown	ret_addr

Execution Lifecycle of a Thread

```
void Thread::dispatch_to(Thread * _thread) {  
    threads_low_switch_to(_thread);  
}
```

push thread_start	eip
push KERNEL_CS	es
push 0	eflags
push t_fun_1	ret_addr
push tshutdown	ret_addr

Execution Lifecycle of a Thread

```
static void thread_start() {  
    /* This function is used to release  
       the thread for execution in the ready queue. */  
    /* We need to add code, but it is probably  
       nothing more than enabling interrupts. */  
}
```

push t_fun_1	ret_addr
push tshutdown	ret_addr

Execution Lifecycle of a Thread

```
void t_fun_1() {  
    /* do_something_in_this_thread */  
}
```

push **tshutdown**

ret_addr

Execution Lifecycle of a Thread

```
static void tshutdown() {  
    /* terminates thread after thread func returns*/  
}
```

Execution Lifecycle of a Thread

Low-Level Dispatching on the x86 : Summary

We learned:

- Dispatching can be **implemented** similarly to **exception handlers**.
(Side Note: This approach is very **general** and **portable**.)
 - Implementation of the **Thread Control Block** in C++ class **Thread**
 - **Dispatching in 4 Steps**: Start, Store, Locate, Load
 - **Thread Creation**
 - The **Execution Lifecycle** of a Thread
-