


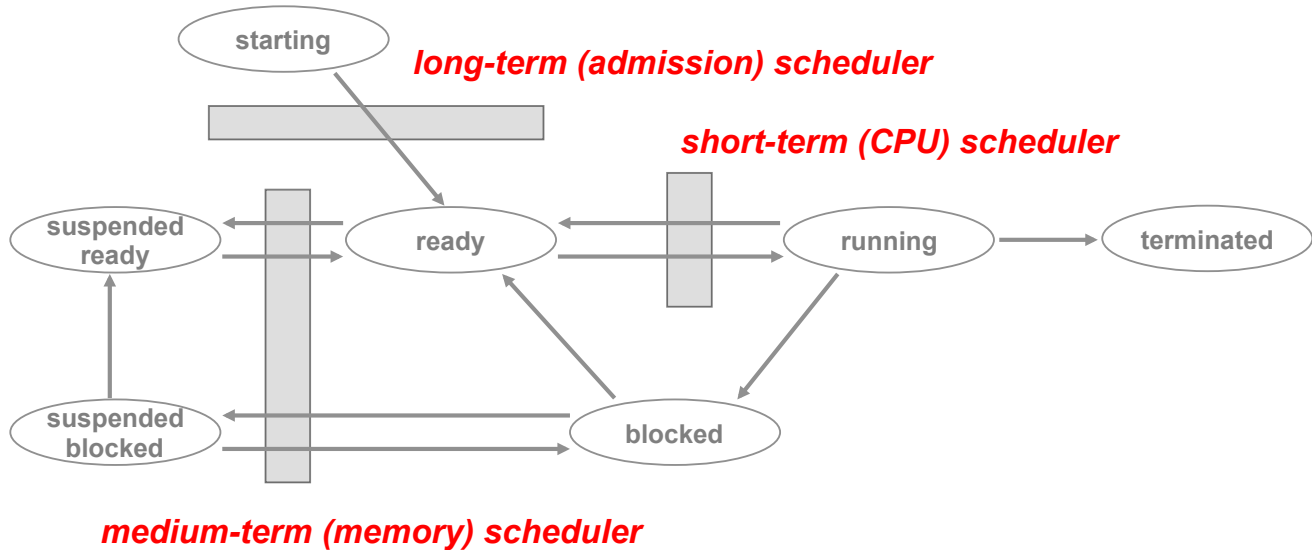
Concurrency and Threading: CPU Scheduling

- Execution State of a Thread or Process
- Schedulers in the OS
- Structure of CPU Schedulers
- Criteria for Scheduling
- Scheduling Algorithms: FIFO, SJF, FP, MLFQ
- Multiprocessor Scheduling


State of a Thread/Process



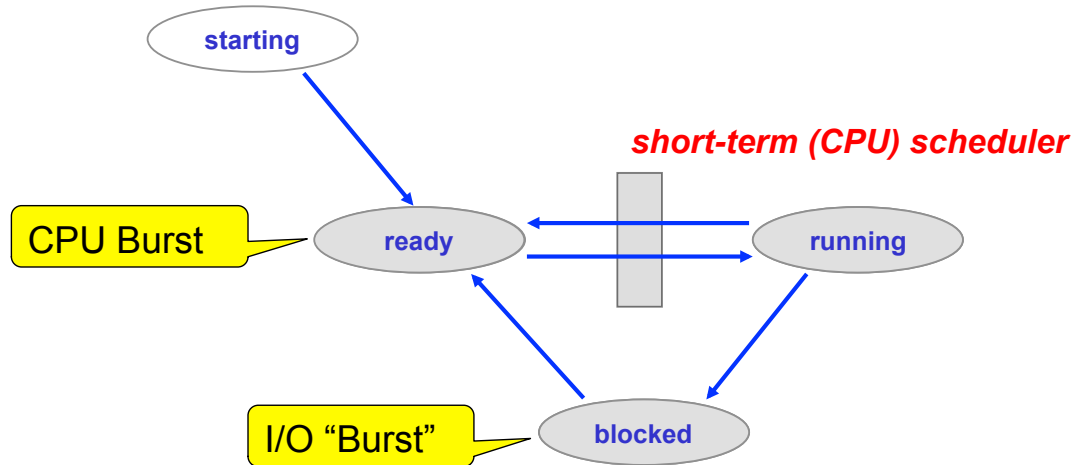
Schedulers




Multiprogramming




CPU Bursts vs. I/O Bursts




CPU Bursts and Jobs



CPU Bursts and Jobs

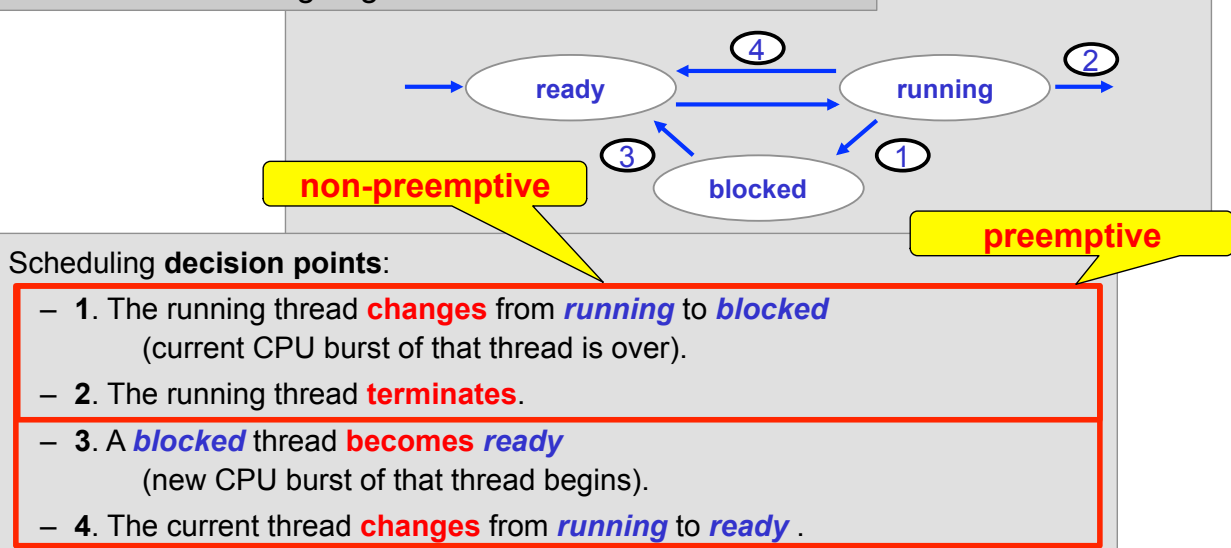


CPU Bursts and Jobs

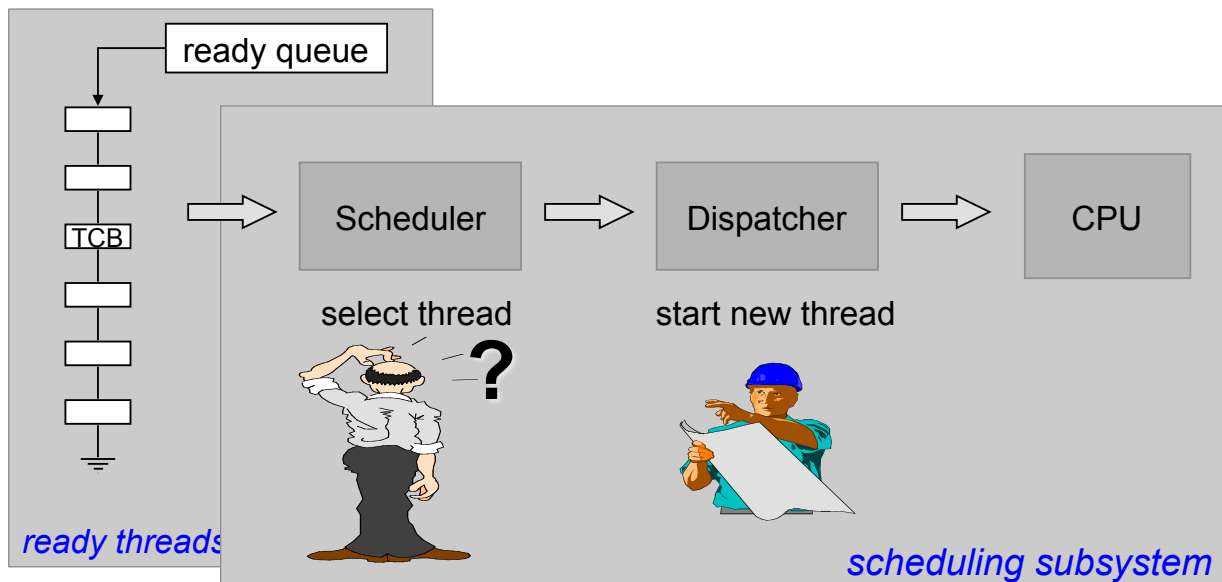


Scheduling Decisions


Decision: “Who is going to use the CPU next?!”



Structure of a Scheduling Subsystem



What Is a Good Scheduler? Criteria



User oriented:

Response time : time interval from start of job to first response

Normalized response time: ratio of response time to execution time


Waiting time : sum of periods spent waiting in ready queue

System oriented:

CPU utilization : percentage of time CPU is busy

Throughput : number of jobs completed per time unit

What Is a Good Scheduler? Criteria



Any good scheduler should:

maximize CPU utilization and throughput

minimize waiting time, response time


Huh?

maximum/minimum values?

average values?

variance?

Scheduling Algorithms




- **FCFS** : First-come-first-served
- **SPN**: Shortest Process Next
- **Priority** scheduling
- **RR** : Round-Robin
- **MLFQ**: Multilevel Feedback Queue Scheduling
- **Multiprocessor** scheduling

Common Structure of a Scheduler



First-Come-First-Served (FCFS/FIFO)



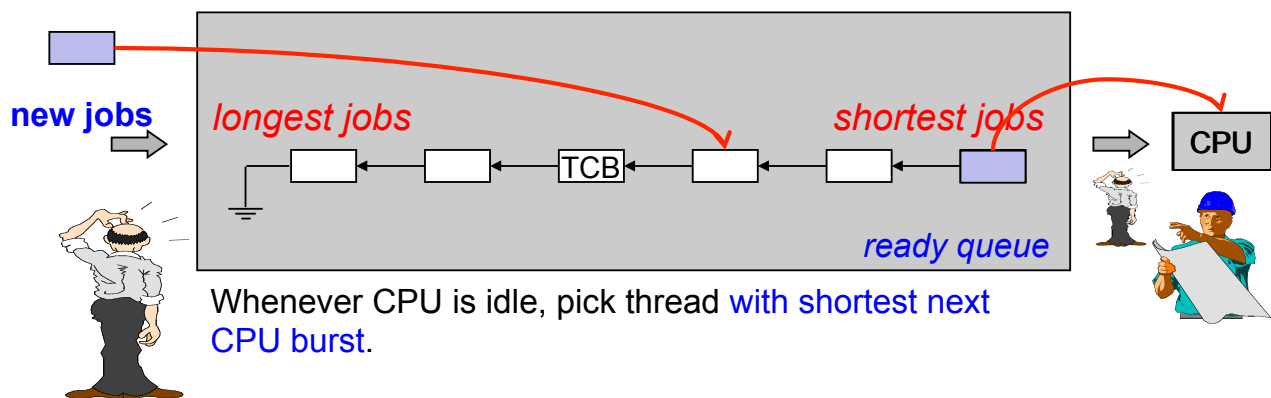
Pros:

- very simple

Cons:

- long average and worst-case waiting times
- poor dynamic behavior (convoy effect)

Shortest-Job-First



Whenever CPU is idle, pick thread with shortest next CPU burst.

Pros:


- minimizes average waiting times

Cons:

- **Starvation** of threads with long CPU bursts.
- How to determine length of next CPU burst?!

SJF Minimizes Average Waiting Time

Provably optimal! Proof: swapping of jobs



$$dW = t_{\text{short}} - t_{\text{long}} < 0$$

Example:

6	12	8	4
6	8	12	4
6	8	4	12
6	4	8	12
4	6	8	12

$$W = 6 + 18 + 26 = 50$$

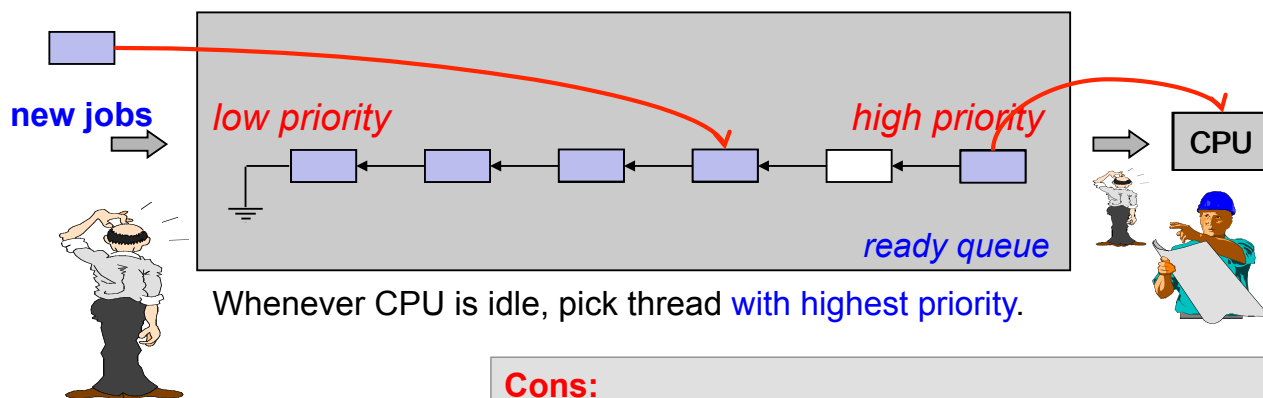
$$W = 6 + 14 + 26 = 46$$

$$W = 6 + 14 + 18 = 38$$

$$W = 6 + 10 + 18 = 34$$

$$W = 4 + 10 + 18 = 32$$

Fixed-Priority Scheduler




Priority:

- thread/process class
- urgency

Cons:

- **Starvation** of low-priority threads

Fixed-Priority Scheduler



Priority:

- thread/process class
- urgency

Cons:


- **Starvation** of low-priority threads

Solution:

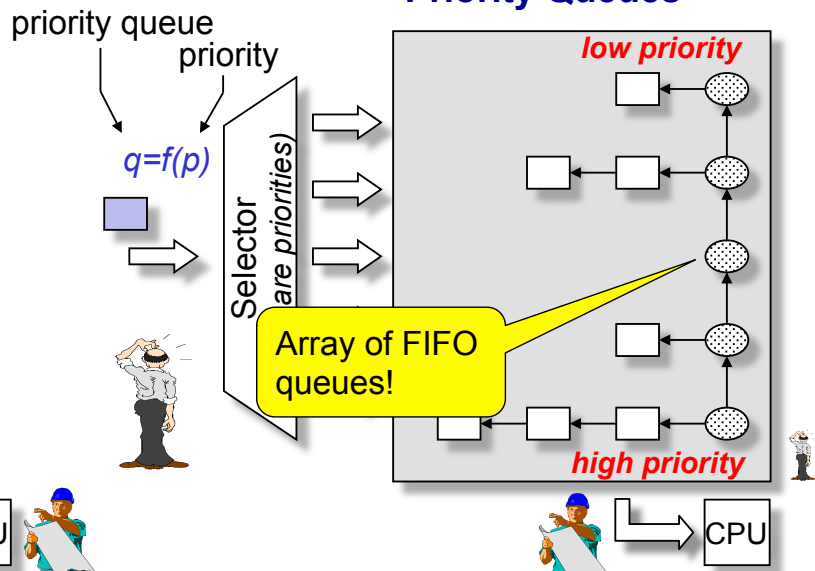
- **Aging**: increase priority over time

Fixed Priority Scheduling (implementation)


Conceptually



Priority Queues



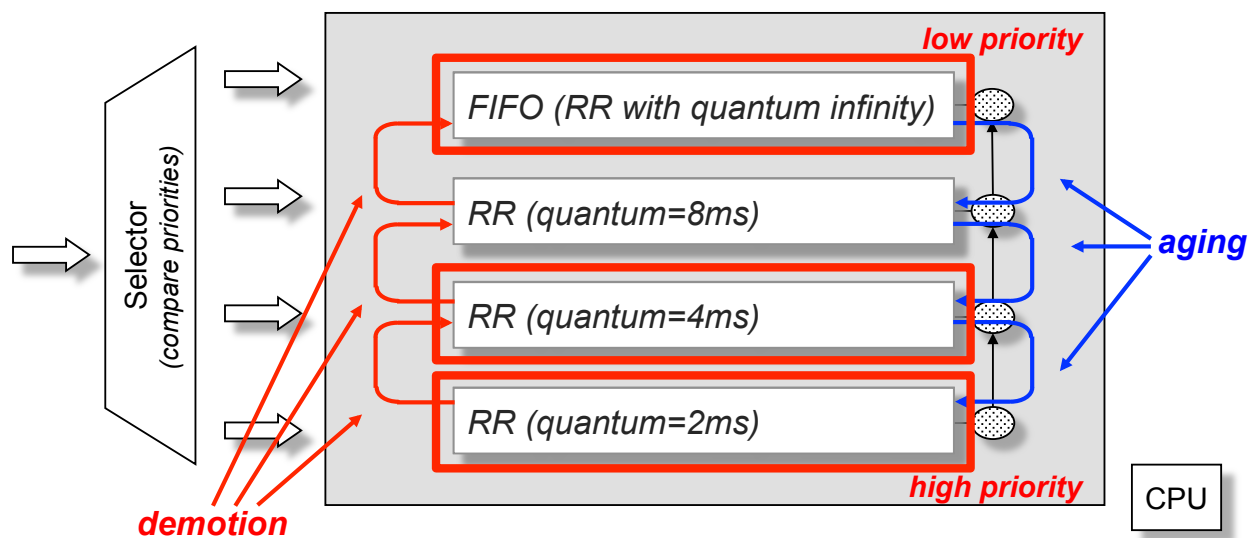
Round-Robin



- FIFO with **preemption** after **time quantum**
- Popular method for **time sharing**
- Choice of **time quantum**:
 - **large**: FCFS
 - **small**: "Processor sharing"
- Time quantum also defines **context-switching overhead**

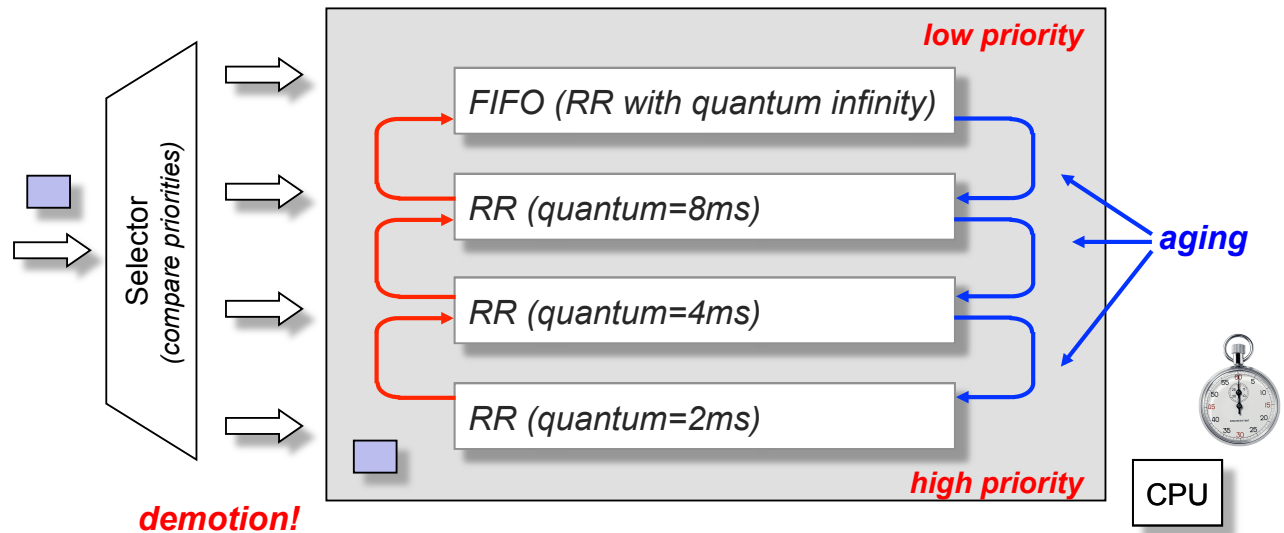
Multilevel Feedback Queue Scheduling

(conceptually!)

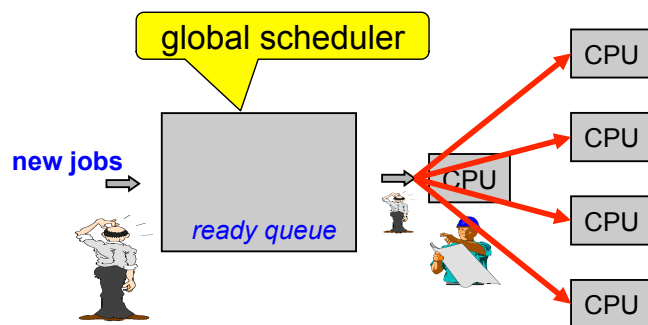


Multilevel Feedback Queue Scheduling

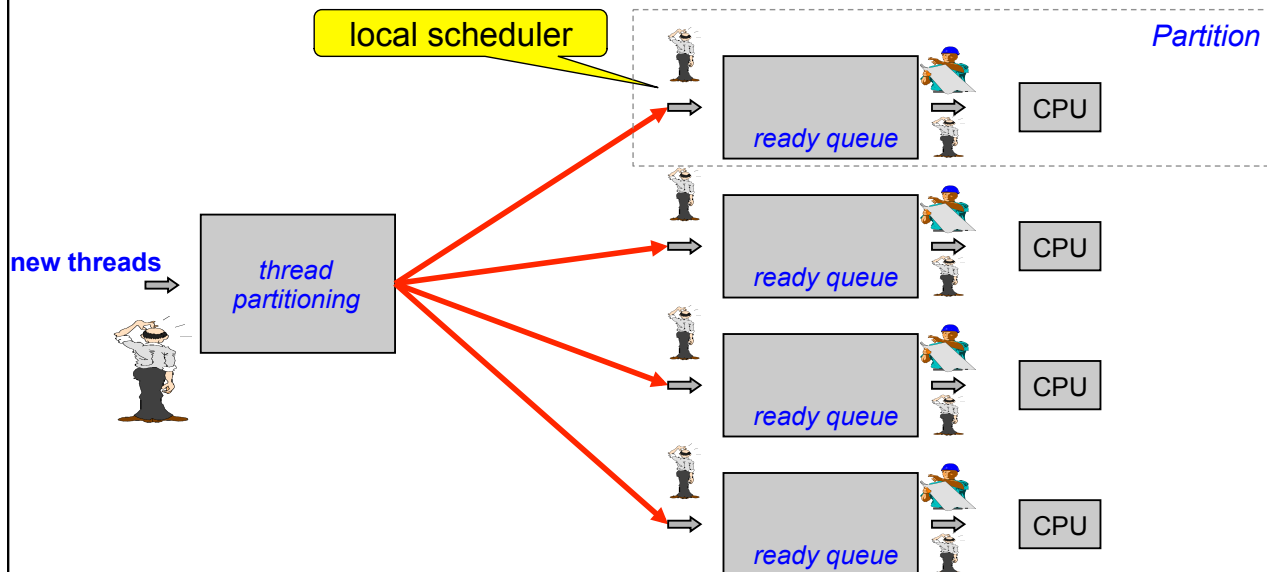
(conceptually!)



Dynamic Multiprocessor Scheduling



Static Multiprocessor Scheduling



CPU Scheduling: Summary

- Execution State of a Thread or Process
- Schedulers in the OS
- Structure of Schedulers
- Criteria for Scheduling
- Scheduling Algorithms: FIFO, SJF, FP, MLFQ
- Multiprocessor Scheduling