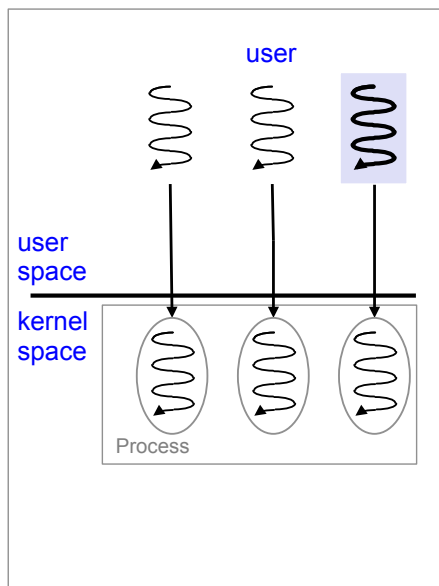


## User-Level vs. Kernel-Level Threads

- Kernel-Level Threads
- User-Level Threads
- Advantages and Limitations of User-Level Threads
- Kernel-level Support for User-Level Threads?
- Example: Scheduler Activations

Thomas E. Anderson, Brian N. Bershad, Edward D. Lazowska, and Henry M. Levy, "Scheduler Activations: Effective Kernel Support for the User-level Management of Parallelism". ACM SIGOPS Operating Systems Review, Volume 25, Issue 5, Oct. 1991.

## Kernel-Level Threads



### Kernel-Level Threads:

- Threads managed by Operating System

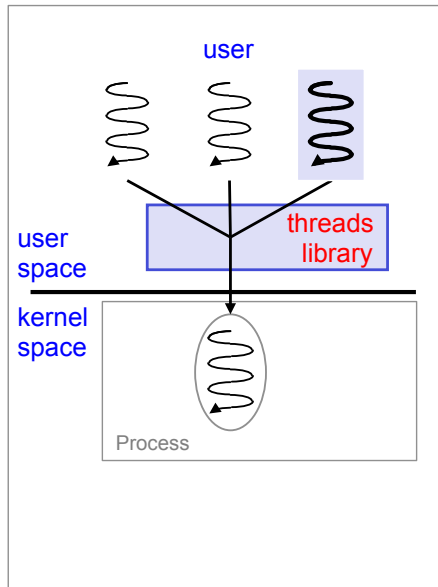
### Pros:

- Avoid system-integration issues (see later)

### Cons:

- Too heavy-weight!

## User-Level Threads



### User-Level Threads:

- Managed by runtime library.
- Management operations require no kernel intervention.

## User-Level Threads using Standard C Library

```
int getcontext(ucontext_t *ucp);
```

saves current thread's execution context in the structure pointed to by [ucp](#).

**ucontext\_t**: structure type suitable for holding the context for a user thread of execution. It contains at least these fields:

<code>mcontext_t uc_mcontext;</code>	saved registers
<code>stack_t uc_stack;</code>	stack area
<code>sigset_t uc_sigmask;</code>	signals being blocked
<code>ucontext_t *uc_link;</code>	context to assume when this one returns

The `uc_link` field points to the context to resume when this context's entry point function returns. If `uc_link` is equal to NULL, then the process exits when this context returns.

## User-Level Threads using Standard C Library

```
int getcontext(ucontext_t *ucp);
```

saves current thread's execution context in the structure pointed to by `ucp`.

```
int setcontext(const ucontext_t * ucp);
```

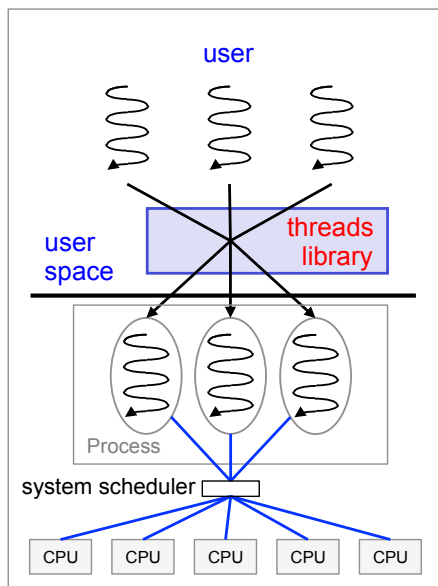
make a previously saved thread context the current thread context. The current context is lost and `setcontext()` does not return.

Instead, execution continue in the context specified by `ucp`.

```
void makecontext(ucontext_t *ucp, void (*func)(), int argc, ...);
```

modifies context pointed to by `ucp` so that it will continue execution by invoking `func()` with the arguments provided. Context `ucp` must have previously been initialized by call to `getcontext()`.

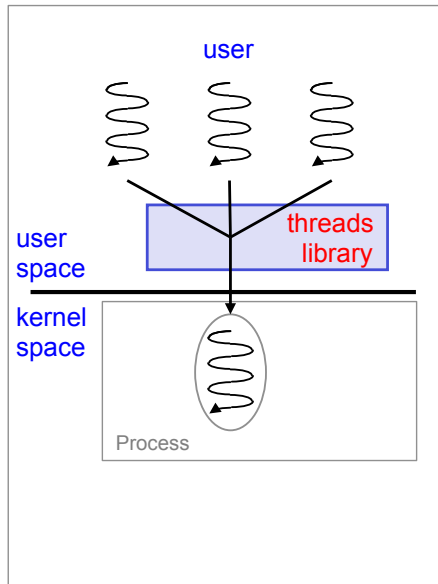
## User-Level Threads: in Practice



### User-Level Threads:

- Managed by runtime library.
- Management operations require no kernel intervention.

## User-Level Threads



### User-Level Threads:

- Managed by runtime library.
- Management operations require no kernel intervention.

### Pros:

1. Low-cost
2. Flexible (various APIs: POSIX, Actors, Java,...)
3. Implementation requires no change to OS.

### Cons:

- Performance issues due to mapping to OS resources (see later)

## User-Level Threads: Advantages

**Kernel-level threads** have inherent **disadvantages**

- **Performance Cost of thread management operations:** Must cross protection boundary on every thread operation, even for operations on threads of the same process.
- **Cost of generality:** A **single implementation** must be used **by all applications**. (In contrast, user-level libraries can be **tuned to applications**.)

Operation	user-level threads	kernel-level threads	kernel-level processes
Null Fork	34	948	11300
Signal-Wait	37	441	1840

Relative latency of thread-operation (Anderson et al.)  
(in comparison, procedure call: 7, exception latency: 19)

## User-Level Threads: Limitations

User-level threads are **difficult to use** and to **integrate with system services**:

1. **Kernel events**, such as processor preemption, I/O blocking, and resumption, are handled by the kernel **invisibly to the user level**.
2. Kernel threads are **scheduled obliviously** with respect to the user-level thread state.

### Scenario:

When a user-level thread **blocks**, the kernel thread **also blocks**.

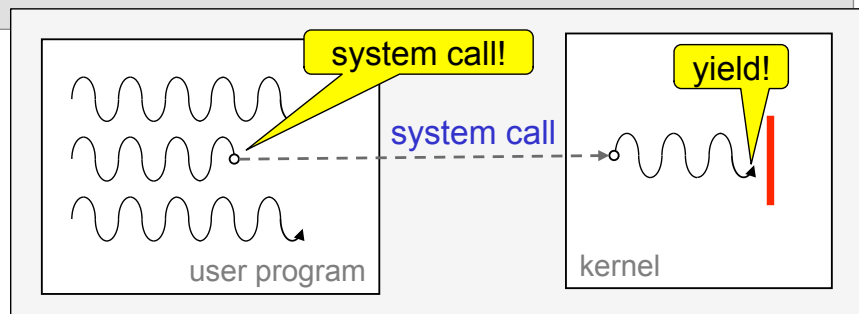
As a result, the physical processor **is lost to the process** while the blocking request is pending.

## User-Level Threads: Limitations

### Scenario:

When a user-level thread **blocks**, the kernel thread serving as its virtual processor **also blocks**.

As a result, the physical processor **is lost to the processes** while the blocking request is pending.



## User-Level Threads: Limitations

### Scenario:

When a user-level thread **blocks**, the kernel thread serving as its virtual processor **also blocks**.

As a result, the physical processor **is lost to the processes** while the blocking request is pending.

**Solution (?)**: Create **more kernel threads**; when one kernel thread blocks because its user-level thread blocks in the kernel, **another kernel thread is available** to run user-level threads on that processor.

**However**: When the thread **unblocks**, there will be **more runnable kernel threads than processors**.

-> OS now decides on behalf of the application which user-level threads to run.

## User-Level Threads: Limitations

**However**: When the thread **unblocks**, there will be **more runnable kernel threads than processors**.

-> OS now decides on behalf of the application which user-level threads to run.

**Solution (?)** : The operating system could employ **some kind of Round Robin** to ensure each thread makes progress.

**However**: When user-level threads are running on top of kernel threads, **round-robin can lead to problems**.

**Example**: A **kernel thread** could be preempted while its **user-level thread** is **holding a spin-lock**;

Any user-level threads accessing the lock will then **spin-wait until the lock holder is re-scheduled**.

## User-Level or Kernel-Level Threads?

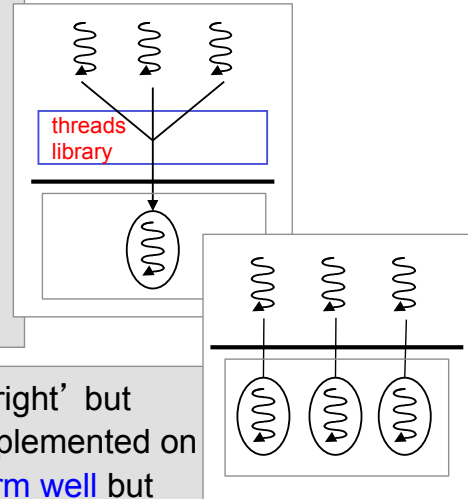
### User-Level Threads:

- (+) Low-cost, flexible, no changes to OS
- (-) Performance issues due to mapping to OS resources

### Kernel-Level Threads:

- (+) Avoid system integration problems
- (-) Too heavyweight

**“Dilemma”:** “Employ **kernel threads**, which ‘work right’ but perform poorly, or employ **user-level threads** implemented on top of kernel threads or processes, which **perform well** but are **functionally deficient**.” (Anderson *et al.*)



## How about Kernel-level Support for ULTs?

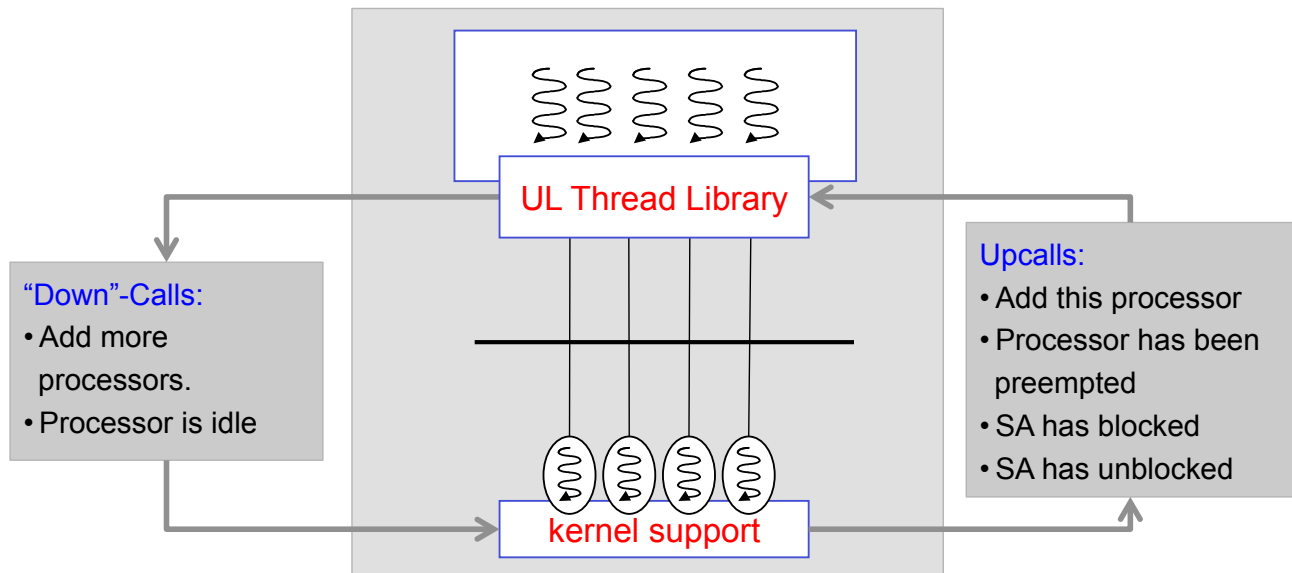
### What if we could ...

- ... **mimic** the behavior of kernel thread management system?
  - i.e., no idling processor in presence of ready threads?
  - i.e., multiprogramming within and across processes?
- ... keep thread management **overhead** same as user-level threads?

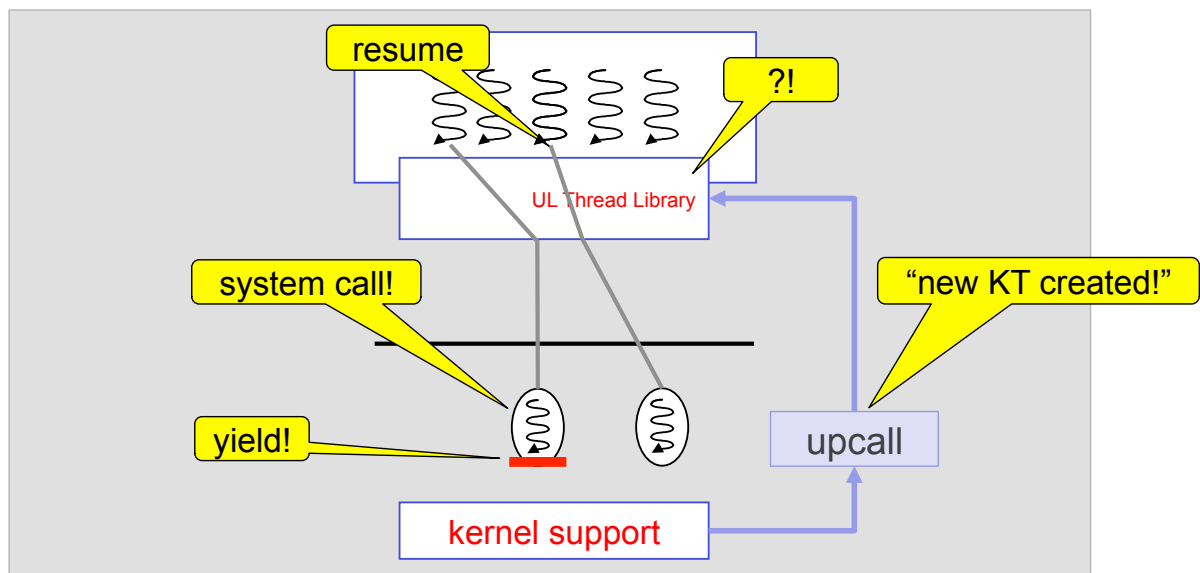
### How about ...

- ... **kernel allocates physical processors** to processes.
- ... UL thread system has complete **control over which thread to run** on allocated processors.
- ... UL thread system knows about suspended/resumed threads in kernel.
- ... UL thread system can request/release processors.

## One Solution: “Scheduler Activations”

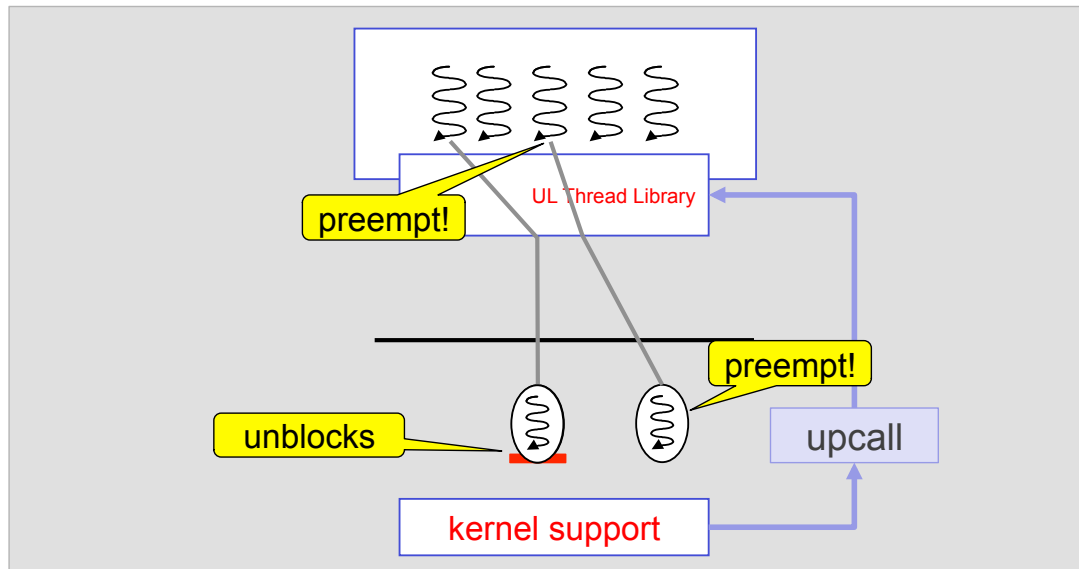


## Handling Blocking Threads using Scheduler Activations





## Resuming Blocked Threads using Scheduler Activations



## User-Level vs. Kernel-Level Threads

- Kernel-Level Threads
- User-Level Threads
- Advantages and Limitations of User-Level Threads
- Kernel-level Support for User-Level Threads?
- Example: Scheduler Activations