

## Event-Based Programming

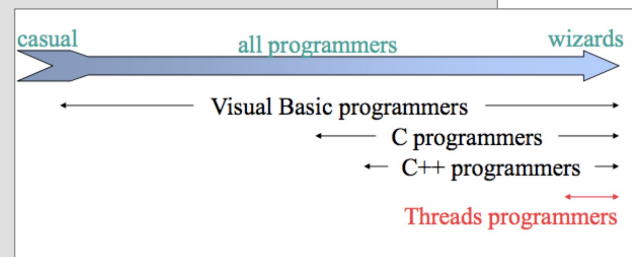
- Threads are “Hard”
- RECAP: Cooperative Task Management
- Event-driven Programming
- Event: Advantages
- Events: Problems

## Threads have Limitations

John Ousterhout on Threads:

**Threads are wrong to use:**

- Threads are “Too hard for most programmers to use.”
- “Even for experts, development is painful.”



John Ousterhout: “Why Threads Are a Bad Idea (for most purposes)”, Sept. 1995

## Threads have Limitations

### John Ousterhout on Threads:

#### Threads are wrong to use:

- Threads are “Too hard for most programmers to use.”
- “Even for experts, development is painful.”

#### Conclusion:

- “Threads should be used only when true CPU concurrency (parallelism) is needed.”
- “For most purposes proposed for threads, events are better”

John Ousterhout: “Why Threads Are a Bad Idea (for most purposes)”, Sept. 1995


## Why Threads are “Hard”

### Synchronization:

- Must synchronize to ensure inter-thread invariants at all times.
- (\*) Forget a lock? Invariant violated. (In English: data corrupted)

### Deadlocks:

- (\*) Circular dependencies among locks.



John Ousterhout: “Why Threads Are a Bad Idea (for most purposes)”, Sept. 1995


## Why Threads are “Hard”

### Synchronization:

- Must **synchronize** to **ensure inter-thread invariants at all times**.
- (\*) **Forget a lock?** Invariant violated. (In English: data corrupted)

### Deadlocks:

- (\*) **Circular dependencies among locks.**

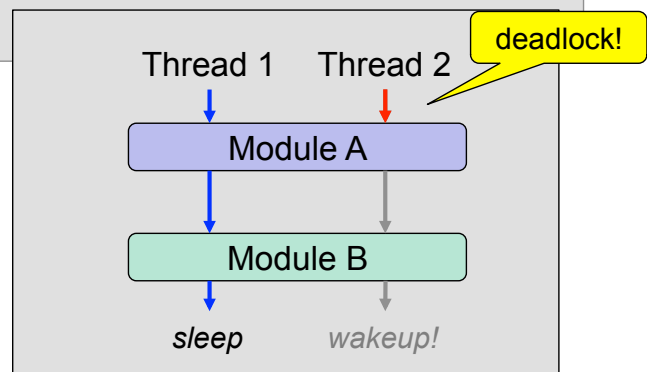


John Ousterhout: “Why Threads Are a Bad Idea (for most purposes)”, Sept. 1995

## Why Threads are “Hard”

### “Lack of Composability”:

- (\*) Cannot **design** modules **independently**.




John Ousterhout: “Why Threads Are a Bad Idea (for most purposes)”, Sept. 1995

## Why Threads are “Hard”

“Locks don’t work with Callbacks”:

(\*)




John Ousterhout: “Why Threads Are a Bad Idea (for most purposes)”, Sept. 1995

## Why Threads are “Hard”

“Achieving good performance is hard”:

- (\*) Simple locking (e.g. monitors) yields low concurrency.
- (\*) Fine-grain locking increases complex and reduces performance in normal case.
- OS limits performance (scheduling, context switches, etc.)

Figure from: M. Welsh, D. Culler, and E. Brewer, SEDA: An Architecture for Well Conditioned, Scalable Internet Services (SOSP, 2001)



John Ousterhout: “Why Threads Are a Bad Idea (for most purposes)”, Sept. 1995


## Why Threads are “Hard”

### Threads “not well supported”:

- Hard to port threaded code.
- Standard libraries often not thread-safe.
- Kernel calls, window systems often not multi-threaded.
- Debugging tools are somewhat lacking.


John Ousterhout: “Why Threads Are a Bad Idea (for most purposes)”, Sept. 1995

## Cooperative Task Management



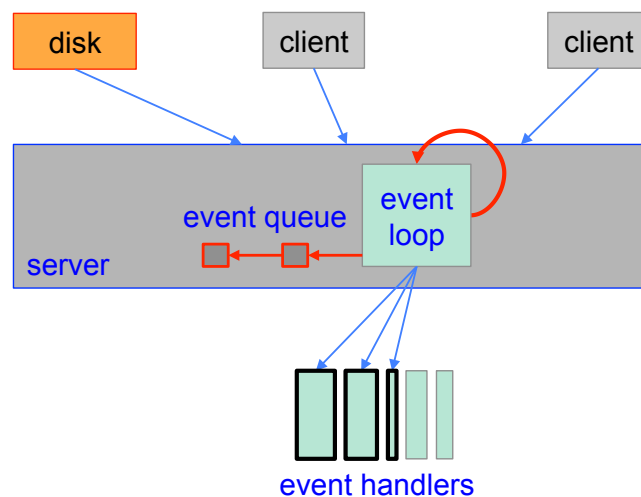
- **Serial** Task Management:
  - Execute each task to completion before starting new task.
- **Cooperative** Task Management:
  - (Voluntarily) yield CPU at well-defined points in execution.
- **Preemptive** Task Management:
  - Execution of tasks can interleave.

## Event-Driven Programming



- Handlers “register” for events
- Event loop **waits for events** and then **invokes handlers** (through function calls)
- Handlers **execute to completion** (no preemption)
- Handlers are typically **short-lived**.
- Events can be initiated by devices as well

## Event-Driven Programming: Step-by-Step



## Example: Event Programming using select/accept

### server

socket()  
bind()  
listen()

select()  
/  
accept()

event handler loop

```
int main() {
    int m_sock = passiveTCPsock(...); /* master socket */
    fd_set rfds, afds;
    FD_ZERO(&afds); FD_SET(m_sock, &afds);
    for (;;) { /* event loop */
        memcpy(&rfds, &afds, sizeof(rfds));
        select(nfds, &rfds, 0, 0, 0);
        if(FD_ISSET(m_sock, &rfds) { /* new connection */
            int s_sock = accept(m_sock, ...); /* "slave" socket */
            FD_SET(s_sock, &afds);
        }
        for(int fd = 0; fd < nfds; fd++)
            if (fd != m_sock && FD_ISSET(fd, &rfds)) {
                EventHandler * handler = identify_handler(fd);
                handler.handle_event(fd); /* invoke event handler */
            }
    }
}
```

## Events: Advantages

### Easier to use:

- No concurrency, no preemption, no synchronization, no deadlocks.

### Easier to debug:

- Timing dependencies all generated by events, not by internal scheduling.
- Problems easier to trace, e.g. slow event handling vs. corrupted memory.

### Fast:

- No context switching, no scheduling overhead.

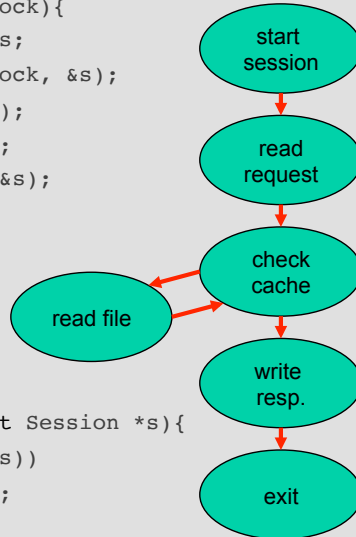
### Portable:

- It's just function calls, after all.

John Ousterhout: "Why Threads Are a Bad Idea (for most purposes)", Sept. 1995

## Events: Manual Control Flow Management

```
thread_main(int sock){
    struct Session s;
    start_session(sock, &s);
    read_request(&s);
    check_cache(&s);
    write_response(&s);
}
```



```
check_cache(struct Session *s){
    if (!in_cache(&s))
        read_file(&s);
}
```

```
StartHandler(event e){
    struct Session * s = new_session(e);
    RequestHandler.enqueue(s);
}

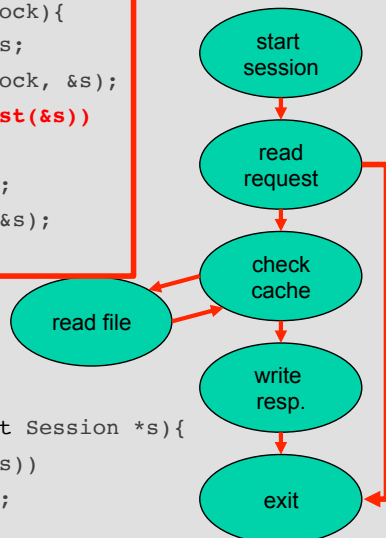
RequestHandler(struct Session *s){
    ...; CacheHandler.enqueue(s);
}

CacheHandler(struct Session *s){
    if (in_cache(&s))
        ResponseHandler.enqueue(s);
    else
        ReadFileHandler.enqueue(s);
}

ExitHandler(struct Session *s){
    ...; free_session(s);
}
```

## Events: Manual Control Flow Management

```
thread_main(int sock){
    struct Session s;
    start_session(sock, &s);
    if (!read_request(&s))
        return;
    check_cache(&s);
    write_response(&s);
}
```



```
check_cache(struct Session *s){
    if (!in_cache(&s))
        read_file(&s);
}
```

```
StartHandler(event e){
    struct Session * s = new_session(e);
    RequestHandler.enqueue(s);
}

RequestHandler(struct Session *s){
    ...; if (error) ExitHandler(s);
    CacheHandler.enqueue(s);
}

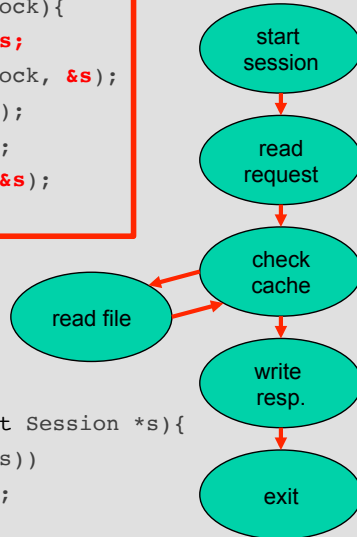
CacheHandler(struct Session *s){
    if (!in_cache(&s))
        ReadFileHandler.enqueue(s);
    else
        ResponseHandler.enqueue(s);
}

ExitHandler(struct Session *s){
    ...; free_session(s);
}
```



## Events: Manual Stack Management

```
thread_main(int sock){
    struct Session s;
    start_session(sock, &s);
    read_request(&s);
    check_cache(&s);
    write_response(&s);
}
```



```
check_cache(struct Session *s){
    if (!in_cache(&s))
        read_file(&s);
}
```

```
StartHandler(event e){
    struct Session * s = new_session(e);
    RequestHandler.enqueue(s);
}

RequestHandler(struct Session *s){
    ...; CacheHandler.enqueue(s);
}

CacheHandler(struct Session *s){
    if (!in_cache(&s))
        ReadFileHandler.enqueue(s);
    else
        ResponseHandler.enqueue(s);
}

ExitHandler(struct Session *s){
    ...; free_session(s);
}
```

## Events: Problems – Summary

### Manual control flow management:

- Control flow for single task is broken up across multiple procedures.

### Manual stack management:

- Have to manually carry local state across these procedures.

### Unexpected blocking:

- What if event handler is blocked by page fault?

### No parallelism:

- How to run events on a multiprocessor?

## Event-Based Programming

---

- Threads are “Hard”
  - RECAP: Cooperative Task Management
  - Event-driven Programming
  - Event: Advantages
  - Events: Problems
-