

UNIX File System Implementation

- Simple implementation of original UNIX File System
 - File System Layout
 - File-System Information and Free-Space Management
 - I-nodes and I-node Management
 - Naming and Directories
 - Evaluation
-

File-System Layout of UNIX FS



File-System Layout of UNIX FS


disk



Boot Block: Contains bootstrap code for operating system.

File-System Layout of UNIX FS

disk




Superblock: Contains information about state of file system, such as size, free space info, etc.

```
struct {
    ushort isize;
    ushort fsize;
    ushort nfree;
    ushort free[100];
    ushort ninode;
    ushort inode[100];
    char flock, ilock, fmod;
    ushort time[2];
}
```

Free Space Management

disk



Superblock: Contains information about state of file system, such as size, free space info, etc.

nfree = 1 **free[..]**

0	0
1	
2	
3	
...	
98	
99	

Free Space Management

disk




Superblock: Contains information about state of file system, such as size, free space info, etc.

nfree = 99 **free[..]**


0	0	
1	FB1	→
2	FB2	→
3	FB3	→
...		→
98	FB98	→
99	FB99	→

Free Space Management

disk



Superblock: Contains information about state of file system, such as size, free space info, etc.




File-System Layout of UNIX FS

disk

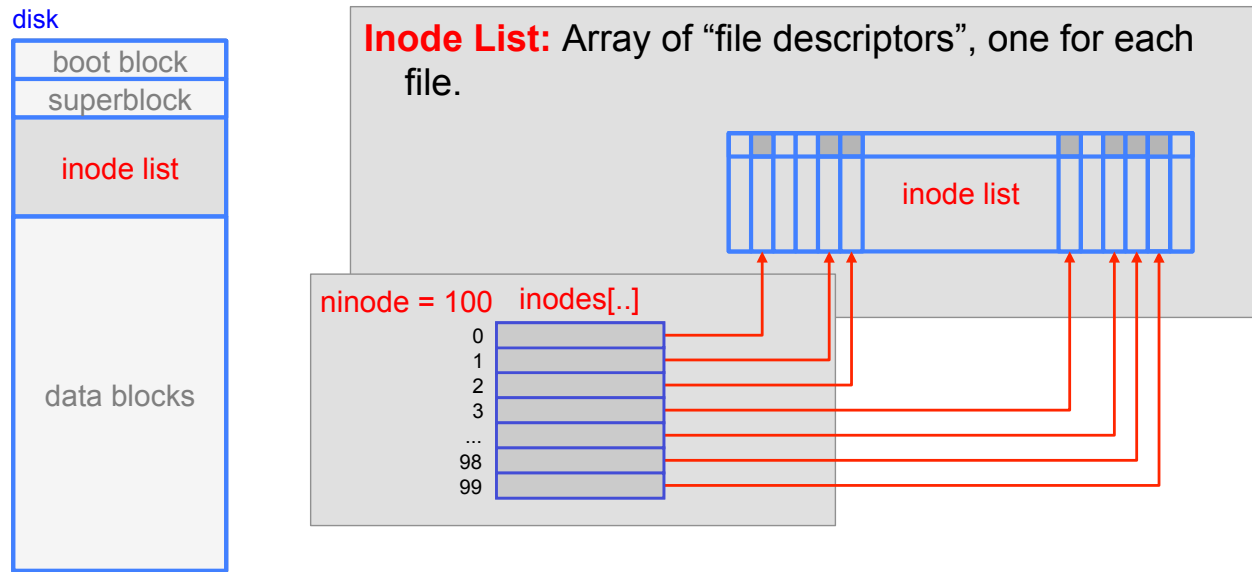


Inode List: Array of meta-data records, one for each file.

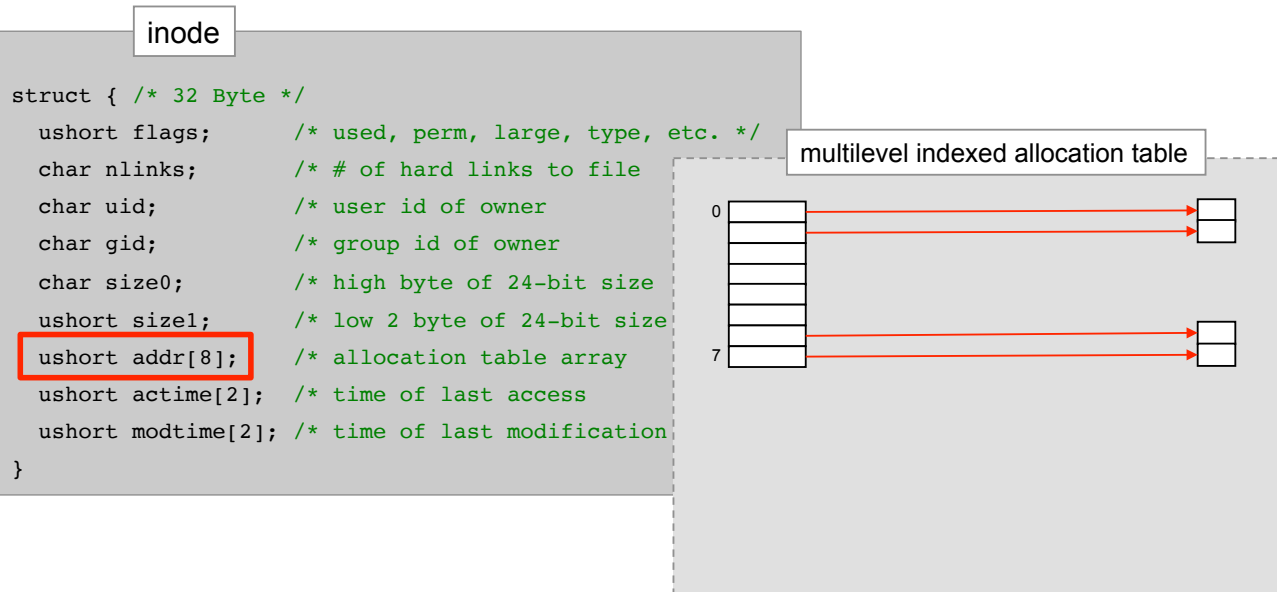
```
struct { /* 32 Byte */
    ushort flags;
    char nlinks;
    char uid;
    char gid;
    char size0;
    ushort size1;
    ushort addr[8];
    ushort actime[2];
    ushort modtime[2];
}
```



File-System Layout of UNIX FS



UNIX inodes



UNIX inodes


inode

```

struct { /* 32 Byte */
    ushort flags;      /* used, perm, large, type, etc. */
    char nlinks;        /* # of hard links to file
    char uid;           /* user id of owner
    char gid;           /* group id of owner
    char size0;         /* high byte of 24-bit size
    ushort size1;       /* low 2 byte of 24-bit size
    ushort addr[8];     /* allocation table array
    ushort actime[2];   /* time of last access
    ushort modtime[2]; /* time of last modification
}

```

multilevel indexed allocation table



UNIX inodes


inode

```

struct { /* 32 Byte */
    ushort flags;      /* used, perm, large, type, etc. */
    char nlinks;        /* # of hard links to file
    char uid;           /* user id of owner
    char gid;           /* group id of owner
    char size0;         /* high byte of 24-bit size
    ushort size1;       /* low 2 byte of 24-bit size
    ushort addr[8];     /* allocation table array
    ushort actime[2];   /* time of last access
    ushort modtime[2]; /* time of last modification
}

```

multilevel indexed allocation table



UNIX inodes

inode


```

struct { /* 32 Byte */
    ushort flags;      /* used, perm, larg
    char nlinks;        /* # of hard links
    char uid;           /* user id of owner
    char gid;           /* group id of owner
    char size0;         /* high byte of 24-
    ushort size1;       /* low 2 byte of 24-
    ushort addr[8];     /* allocation table
    ushort actime[2];   /* time of last acc
    ushort modtime[2]; /* time of last modification */
}

```

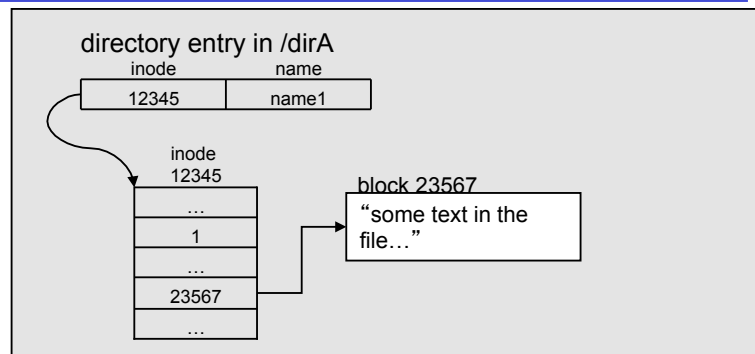
Name information is contained in separate **Directory Files**, which contain entries of type:

(inode number of file, name of file)

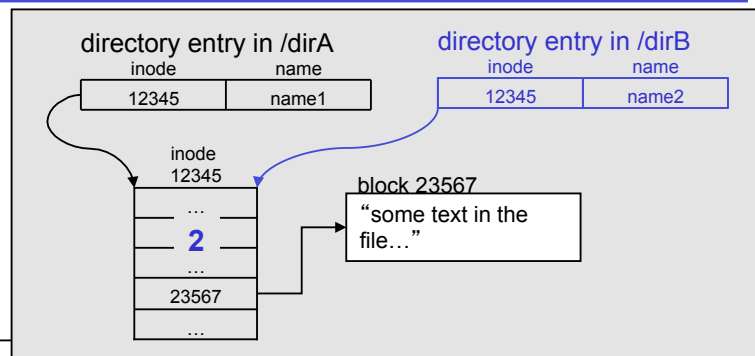


where is the file name?!

Files can have multiple Names (Hard Links)



Files can have multiple Names (Hard Links)

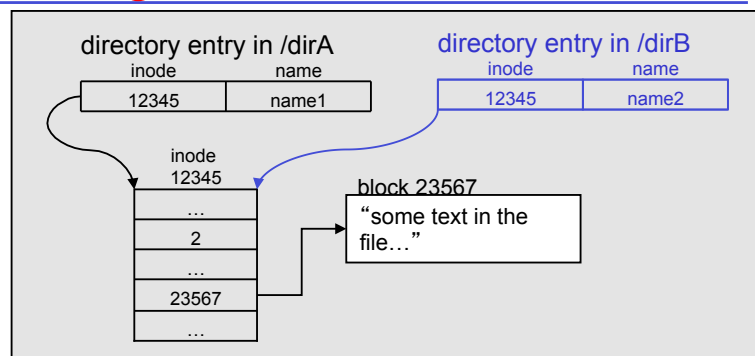


The **link** system call:

```
#include <stdio.h>
#include <unistd.h>

if (link("/dirA/name1", "/dirB/name2") == -1)
    perror("failed to make new link in /dirB");
```

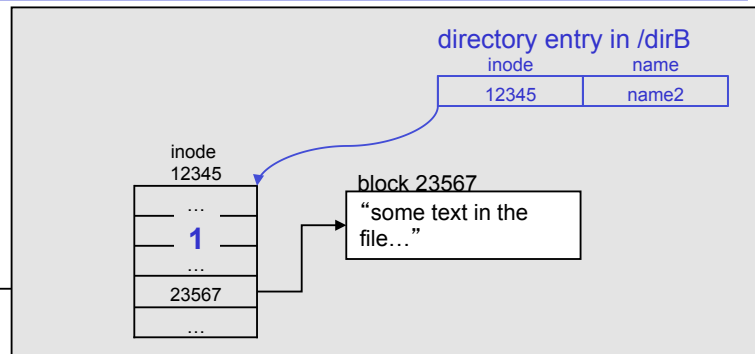
Deleting Files



Deleting Files

```
#include <stdio.h>
#include<unistd.h>

if (unlink("/dirA/name1") == -1)
    perror("failed to delete link in /dirA");
```

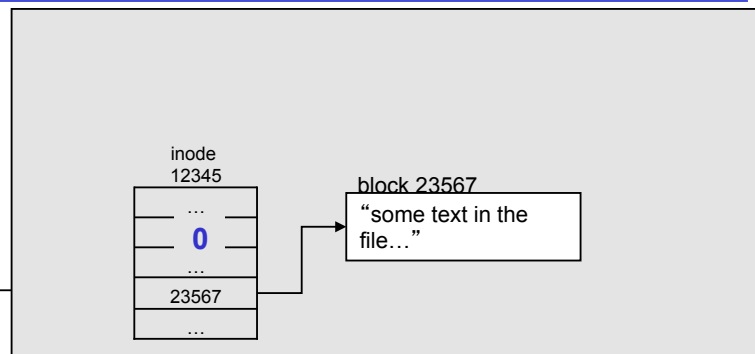


Deleting Files

```
#include <stdio.h>
#include<unistd.h>

if (unlink("/dirA/name1") == -1)
    perror("failed to delete link in /dirA");

if (unlink("/dirB/name2") == -1)
    perror("failed to delete link in /dirB");
```

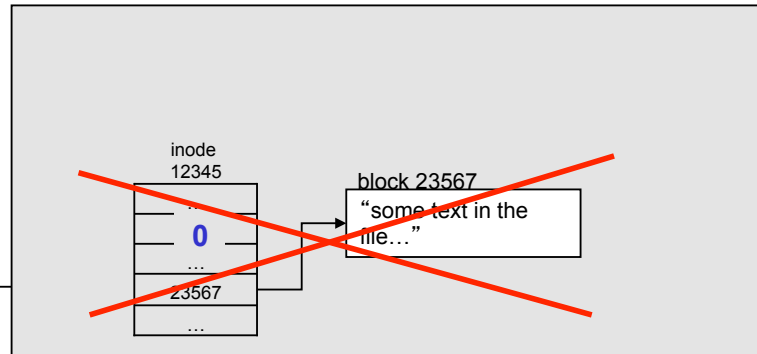


Deleting Files


```
#include <stdio.h>
#include <unistd.h>

if (unlink("/dirA/name1") == -1)
    perror("failed to delete link in /dirA");


if (unlink("/dirB/name2") == -1)
    perror("failed to delete link in /dirB");
```



File-System Layout of UNIX FS



Inode List: Array of “file descriptors”, one for each file.



Pros/Cons of Original Unix File System

Pros:

- Data in small files can be accessed directly from the inode.
 - One read to fetch inode, another to fetch first data block.
- Larger files can be accessed efficiently
- No external fragmentation.

Cons:

- Original file system uses 512-byte blocks.
- Inodes kept separately from data, causing long seeks to access data.
- Inodes of files in a common directory not kept together, causing low performance when searching directories.
- Data blocks of a file are not stored together.
- Free list quickly scrambles, increasing overhead of finding free blocks.

Pros/Cons of Original Unix File System

Pros:

- Data in small files can be accessed directly from the inode.
 - One read to fetch inode, another to fetch first data block.
- Larger files can be accessed efficiently
- No ext

Measurements of original file system used as little as 2% of the available disk bandwidth!

Cons:

- Original file system uses 512-byte blocks.
- Inodes kept separately from data, causing long seeks to access data.
- Inodes of files in a common directory not kept together, causing low performance when searching directories.
- Data blocks of a file are not stored together.
- Free list quickly scrambles, increasing overhead of finding free blocks.

UNIX File System Implementation

- Simple implementation of original UNIX File System
 - File System Layout
 - Free-Space Management
 - I-nodes and I-node Management
 - Evaluation
-