

The Fast File System

M.K. McKusick, W. Joy, S.L. Leffler, R.S. Fabry, "A Fast File System for UNIX",
ACM Transactions on Computer Systems, Vol. 2, No. 3, Aug. 1984

- Recall: Limitations of Original UNIX File System
- Fast File System (FFS)
 - Increase Block Size
 - Make File System "disk-aware"
 - Some Functional Improvements

Recall: Problems with Original Unix File System

- Original file system uses 512/1024-byte blocks.
- Inodes kept separately from data, causing long seeks to access data.
- Inodes of files in a common directory not kept together, causing low performance when searching directories.
- Data blocks of a file are not stored together.
- Free list quickly scrambles, increasing overhead of finding free blocks.

Original file system treats disk as Random Access Device!

“Fast FS” (FFS, ca. 1984)

M.K. McKusick, W. Joy, S.L. Leffler, R.S. Fabry, “A Fast File System for UNIX”,
ACM Transactions on Computer Systems, Vol. 2, No. 3, Aug. 1984

Three-pronged approach:

1. Increase **block size**
2. Make file system **disk-aware**
3. **Functional** improvements

FFS: Increase Block Size

Increase **block size** from **512 byte** to at least **4096 byte**.

- Block size **power-of-two multiple of 4096 byte**
- Defined **during file system creation**
- Stored as **parameter** in superblock

FFS: Block Size and Fragmentation

Observation: Typical UNIX files are **small** compared to large blocks.

M.K. McKusick, *et al.*


Table I. Amount of Wasted Space as a Function of Block Size

Space used (Mbytes)	% Waste	Organization
775.2	0.0	Data only, no separation between files
807.8	4.2	Data only, each file starts on 512-byte boundary
828.7	6.9	Data + inodes, 512-byte block UNIX file system
866.5	11.8	Data + inodes, 1024-byte block UNIX file system
948.5	22.4	Data + inodes, 2048-byte block UNIX file system
1128.3	45.6	Data + inodes, 4096-byte block UNIX file system

Blocks vs. Fragments

Solution: Split blocks into **Fragments**

- **Fragment size** defined at file system creation time
- Example: 4096/1024 has 4kB blocks and 1kB fragments.




FFS Organization: Making FS Disk-Aware

1. Cylinder Groups

2. File System Parameterization

FFS Organization: Cylinder Groups



Cylinder Groups

- groups of multiple adjacent disk cylinders.
- each group maintains own copy of superblock, inode bitmap, data bitmap, inodes, and data blocks:



Allocation of directories and files:

- “keep related stuff together”
- inodes and data blocks
- blocks of same file
- files and directories

FFS: Minimizing Rotational Latency

File System Parameterization


Goal: Parameterize **processor capabilities** and **disk characteristics** so that blocks can be **allocated** in an optimal, configuration-dependent way.

Allocate new block **rotationally well-positioned**.

Disk Parameters:

number of blocks per **track**

disk **spin rate**.



CPU Parameters:

expected time to **service interrupt** and **schedule new disk transfer**

FFS: Performance

Table IIa. Reading Rates of the Old and New UNIX File Systems

Type of file system	Processor and bus measured	Speed (Kbytes/s)	Read bandwidth %	% CPU
Old 1024	750/UNIBUS	29	29/983 3	11
New 4096/1024	750/UNIBUS	221	221/983 22	43
New 8192/1024	750/UNIBUS	233	233/983 24	29
New 4096/1024	750/MASSBUS	466	466/983 47	73
New 8192/1024	750/MASSBUS	466	466/983 47	54

Table IIb. Writing Rates of the Old and New UNIX File Systems

Type of file system	Processor and bus measured	Speed (Kbytes/s)	Write bandwidth %	% CPU
Old 1024	750/UNIBUS	48	48/983 5	29
New 4096/1024	750/UNIBUS	142	142/983 14	43
New 8192/1024	750/UNIBUS	215	215/983 22	46
New 4096/1024	750/MASSBUS	323	323/983 33	94
New 8192/1024	750/MASSBUS	466	466/983 47	95

M.K. McKusick, *et al.*

FFS: Functional Enhancements

1. Arbitrary-length File Names
 2. File Locking
 3. Symbolic Links
 4. “Rename” system call added
 5. Quotas
-

The Fast File System

- Recall: Limitations of Original UNIX File System
 - Fast File System (FFS)
 - Increase Block Size
 - Make File System “disk-aware”
 - Some Functional Improvements
-